

deep_learning_flights

November 18, 2020

```
[1]: from numpy import loadtxt
import tensorflow as tf
from tensorflow import keras
import pandas as pd
from matplotlib import pyplot as plt
import numpy as np

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error
```

```
[2]: # Read in and shuffle dataset
dataset = pd.read_csv("../data_raw/full_data_one_hot.csv").sample(frac=1)

print(dataset.shape)

dataset
```

(581328, 203)

```
[2]:
```

	departure_delay	time_taxi_out	elapsed_time_scheduled	delay_carrier	\
50738	-4	16	61	0	
195908	13	21	145	0	
462546	43	14	102	0	
49188	92	12	126	92	
318575	-10	13	213	0	
...	
364474	1	12	283	0	
88432	-4	19	151	0	
483395	43	16	100	0	
207499	-6	22	235	0	
306795	-3	10	70	0	

	delay_weather	delay_nas	delay_security	delay_late_aircraft	\
50738	0	0	0	0	
195908	13	6	0	0	
462546	0	0	0	22	
49188	0	7	0	0	
318575	0	0	0	0	

...
364474	0	0	0	0
88432	0	0	0	0
483395	0	10	0	20
207499	0	0	0	0
306795	0	0	0	0

	humidity	dew_point	...	wind_direction_NW	wind_direction_S	\
50738	79	69	...	0	0	
195908	85	25	...	0	0	
462546	57	57	...	0	0	
49188	85	75	...	0	0	
318575	14	27	...	0	0	

...
364474	39	27	...	0
88432	73	63	...	0
483395	70	51	...	1
207499	46	54	...	1
306795	51	48	...	0

	wind_direction_SE	wind_direction_SSE	wind_direction_SSW	\
50738	0	1	0	
195908	0	0	0	
462546	0	0	0	
49188	0	0	0	
318575	0	0	0	

...
364474	0	0	0
88432	0	0	0
483395	0	0	0
207499	0	0	0
306795	0	1	0

	wind_direction_SW	wind_direction_VAR	wind_direction_W	\
50738	0	0	0	
195908	0	0	0	
462546	0	0	0	
49188	1	0	0	
318575	0	0	0	

...
364474	0	0	1
88432	0	0	0
483395	0	0	0
207499	0	0	0
306795	0	0	0

wind_direction_WNW	wind_direction_WSW
--------------------	--------------------

50738	0	0
195908	0	0
462546	1	0
49188	0	0
318575	0	0
...
364474	0	0
88432	0	0
483395	0	0
207499	0	0
306795	0	0

[581328 rows x 203 columns]

```
[3]: # Get training labels
Y = dataset.head(400000)['time_taxi_out'].values
Y
```

```
[3]: array([16, 21, 14, ..., 15, 15, 13])
```

```
[4]: # Get training features
X = dataset.head(400000).drop('time_taxi_out', axis=1).values
X.shape
```

```
[4]: (400000, 202)
```

1 Deep Learning

```
[5]: # Create and fit model
model = keras.Sequential()

model.add(keras.layers.Dense(100, input_dim = X.shape[1], activation = 'sigmoid'))
model.add(keras.layers.Dense(30, activation = 'sigmoid'))
model.add(keras.layers.Dense(10, activation = 'relu'))
model.add(keras.layers.Dense(1, activation = 'relu'))

model.compile(loss = 'mse', optimizer = 'adam', metrics = [keras.metrics.
    MeanAbsoluteError()])

history = model.fit(X, Y, validation_split = 0.2, epochs = 10, batch_size = 200)
```

```
Epoch 1/10
1600/1600 [=====] - 2s 1ms/step - loss: 211.6716 -
mean_absolute_error: 9.7036 - val_loss: 158.6571 - val_mean_absolute_error:
8.3450
```

```


Epoch 2/10
1600/1600 [=====] - 2s 2ms/step - loss: 103.7929 -
mean_absolute_error: 6.5229 - val_loss: 81.7991 - val_mean_absolute_error:
6.0273
Epoch 3/10
1600/1600 [=====] - 2s 2ms/step - loss: 76.1877 -
mean_absolute_error: 5.9604 - val_loss: 73.2011 - val_mean_absolute_error:
5.8295
Epoch 4/10
1600/1600 [=====] - 2s 1ms/step - loss: 71.2044 -
mean_absolute_error: 5.8603 - val_loss: 70.4624 - val_mean_absolute_error:
5.7370
Epoch 5/10
1600/1600 [=====] - 2s 1ms/step - loss: 69.2120 -
mean_absolute_error: 5.8076 - val_loss: 68.1306 - val_mean_absolute_error:
5.7780
Epoch 6/10
1600/1600 [=====] - 2s 1ms/step - loss: 68.3088 -
mean_absolute_error: 5.7848 - val_loss: 68.6899 - val_mean_absolute_error:
5.7697
Epoch 7/10
1600/1600 [=====] - 2s 1ms/step - loss: 67.2939 -
mean_absolute_error: 5.7516 - val_loss: 67.6030 - val_mean_absolute_error:
5.8512
Epoch 8/10
1600/1600 [=====] - 2s 1ms/step - loss: 66.8326 -
mean_absolute_error: 5.7382 - val_loss: 67.4560 - val_mean_absolute_error:
5.8358
Epoch 9/10
1600/1600 [=====] - 2s 1ms/step - loss: 66.4438 -
mean_absolute_error: 5.7329 - val_loss: 66.8489 - val_mean_absolute_error:
5.6814
Epoch 10/10
1600/1600 [=====] - 2s 1ms/step - loss: 66.1275 -
mean_absolute_error: 5.7182 - val_loss: 67.4736 - val_mean_absolute_error:
5.6583

```

```

[6]: # Plot training/validation loss
plt.plot(history.history['mean_absolute_error'])
plt.plot(history.history['val_mean_absolute_error'])
plt.legend(['train', 'val'], loc = 'upper right')
plt.xlabel("epoch")
plt.ylabel("MAE")
plt.show()

```



```
[5]: # Predict values for unseen data
```


```
x_predict = dataset[400000:500000].drop('time_taxi_out', axis=1).values  
y_true = dataset[400000:500000]['time_taxi_out'].values
```

```
[8]: y_predicted = model.predict(x_predict)  
print("MAE training: ", mean_absolute_error(Y, model.predict(X)))  
print("MAE test: ", mean_absolute_error(y_true, y_predicted))
```

MAE training: 5.6347700267148015

MAE test: 5.645663092269897

```
[9]: plt.scatter(y_true, y_predicted, s = 1)  
plt.xlabel("y_true")  
plt.ylabel("y_predicted")  
plt.show()
```



2 Linear Regression

```
[6]: from sklearn.linear_model import LinearRegression
     from sklearn.metrics import mean_absolute_error
```

```
[7]: model = LinearRegression().fit(X, Y)
```

```
[8]: print("r_sq training: ", model.score(X, Y))
     print("r_sq test: ", model.score(x_predict, y_true))
```

```
r_sq training:  0.4238081919600284
```

```
r_sq test:  0.4140695987558387
```

```
[9]: y_predicted = model.predict(x_predict)
     test_mae = mean_absolute_error(y_true, y_predicted)
     print("MAE training: ", mean_absolute_error(Y, model.predict(X)))
     print("MAE test: ", test_mae)
```

```
MAE training:  6.6418192713803546
```

```
MAE test:  6.653655706379016
```

3 Polynomial Regression

```
[6]: from sklearn.preprocessing import PolynomialFeatures
```

```
[7]: x_numeric = X[:, 0:13]
     x_one_hot = X[:, 13:]
     transformer = PolynomialFeatures(degree = 2, include_bias = True)
     transformer.fit(x_numeric)
     x_numeric_ = transformer.transform(x_numeric)
     X_ = np.concatenate((x_numeric_, x_one_hot), axis = 1)
```

```
[8]: model = LinearRegression().fit(X_, Y)
```

```
[9]: x_predict_numeric_ = transformer.transform(x_predict[:, 0:13])
     x_predict_ = np.concatenate((x_predict_numeric_, x_predict[:, 13:]), axis = 1)
     print("r_sq training: ", model.score(X_, Y))
     print("r_sq test: ", model.score(x_predict_, y_true))
```

```
r_sq training:  0.4962595935412766
r_sq test:  0.4982515373554338
```

```
[10]: y_predicted = model.predict(x_predict_)
     test_mae = mean_absolute_error(y_true, y_predicted)
     print("MAE training: ", mean_absolute_error(Y, model.predict(X_)))
     print("MAE test: ", test_mae)
```

```
MAE training:  6.276235768111072
MAE test:  6.2347969036412145
```

```
[ ]:
```